

DEEP LEARNING

Instructure: Dr Mayank Vatsa



Submitted by: Shrikesh Jalan (M21MA012)

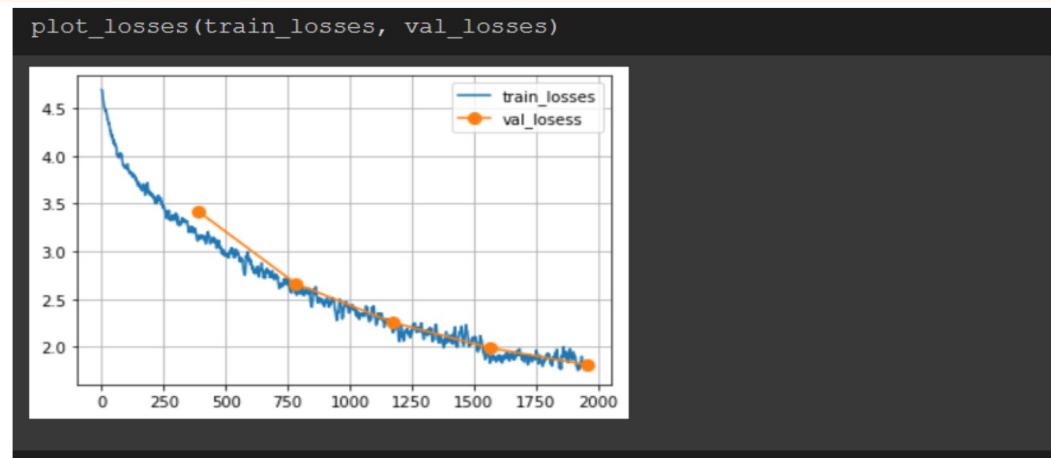
Searching for MobileNetV3 - Results

Dataset: MNIST Fashion

The model is trained on MobileNetV1, MobileNetV2, and MobileNetV3 are trained with fashion-MNIST dataset and the output is given below.

```
trainset = torchvision.datasets.FashionMNIST(root = "./data", train = True, download = True)
testset = torchvision.datasets.FashionMNIST(root = "./data", train = False, download = True)

Extracting http://fashion-mnist.s3-eu-central-1.amazonaws.com/train-idx3-ubyte.gz to ./data/FashionMNIST/raw/train-images-idx3-ubyte.gz
HBox(children=[FloatProgress(value=1.0, bar_style='info', max=1.0, HTML(value=''))])
Extracting http://fashion-mnist.s3-eu-central-1.amazonaws.com/train-labels-idx3-ubyte.gz to ./data/FashionMNIST/raw/train-labels-idx3-ubyte.gz
HBox(children=[FloatProgress(value=1.0, bar_style='info', max=1.0, HTML(value=''))])
Extracting http://fashion-mnist.s3-eu-central-1.amazonaws.com/t10k-idx3-ubyte.gz to ./data/FashionMNIST/raw/t10k-images-idx3-ubyte.gz
HBox(children=[FloatProgress(value=1.0, bar_style='info', max=1.0, HTML(value=''))])
Extracting http://fashion-mnist.s3-eu-central-1.amazonaws.com/t10k-labels-idx3-ubyte.gz to ./data/FashionMNIST/raw/t10k-labels-idx3-ubyte.gz
HBox(children=[FloatProgress(value=1.0, bar_style='info', max=1.0, HTML(value=''))])
Extracting http://fashion-mnist.s3-eu-central-1.amazonaws.com/t10k-images-idx3-ubyte.gz to ./data/FashionMNIST/raw/t10k-labels-idx3-ubyte.gz
HBox(children=[FloatProgress(value=1.0, bar_style='info', max=1.0, HTML(value=''))])
Extracting http://fashion-mnist.s3-eu-central-1.amazonaws.com/t10k-labels-idx3-ubyte.gz to ./data/FashionMNIST/raw/t10k-labels-idx3-ubyte.gz
HBox(children=[FloatProgress(value=1.0, bar_style='info', max=1.0, HTML(value=''))])
Processing...
/usr/local/lib/python3.6/dist-packages/torchvision/datasets/mnist.py:460: UserWarning: The given NumPy array is not writeable, and PyTorch does not support
return torch.from_numpy(numpy_array).copy_(label_view)
MNIST FASHION data download
```



Training of MobileNetV1, MobileNetV2, and MobileNetV3:

```
model = MobileNetV1().to(device)
epochs = 5
lr = 1e-3
train_losses, val_losses = train(train_loader, val_loader, model, epochs, lr)
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
EPOCH # 1
train_loss: 3.093 | train_acc: 0.3165
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
val_loss: 2.415 | val_acc: 0.1584
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
EPOCH # 2
train_loss: 2.907 | train_acc: 0.2468
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
val_loss: 2.657 | val_acc: 0.1582
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
EPOCH # 3
train_loss: 2.437 | train_acc: 0.3630
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
val_loss: 2.395 | val_acc: 0.1584
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
EPOCH # 4
train_loss: 2.333 | train_acc: 0.3930
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
val_loss: 2.333 | val_acc: 0.1584
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
EPOCH # 5
train_loss: 2.197 | train_acc: 0.4268
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
val_loss: 2.197 | val_acc: 0.1584
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
MobileNet training
```

```
start = time.time()
val_acc = get_acc(model, val_loader)
end = time.time()

print('Batch_size={}, epochs={}, lr={}'.format(128, epochs, lr))
print("Val accuracy =", val_acc)
print("Process validation time: {:.4f} s".format(end - start))

HBox(children=[FloatProgress(value=0.0, max=79.0), HTML(value='')])

Batch_size=128, epochs=5, lr=0.001
Val accuracy = 0.5061
Process validation time: 16.1220 s
```

MobileNet1 Validation

```
model = MobileNetV2().to(device)
epochs = 5
lr = 1e-3
train_losses, val_losses = train(train_loader, val_loader, model, epochs, lr)
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
EPOCH # 1
train_loss: 3.707 | train_acc: 0.1116
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
val_loss: 3.713 | val_acc: 0.1397
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
EPOCH # 2
train_loss: 3.057 | train_acc: 0.2223
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
val_loss: 2.7629 | val_acc: 0.2014
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
EPOCH # 3
train_loss: 2.6399 | train_acc: 0.3165
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
val_loss: 2.6399 | val_acc: 0.3595
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
EPOCH # 4
train_loss: 2.4011 | train_acc: 0.3931
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
val_loss: 2.4011 | val_acc: 0.3595
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
EPOCH # 5
train_loss: 2.1111 | train_acc: 0.4268
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
val_loss: 2.1111 | val_acc: 0.3595
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
MobileNet2 Training
```

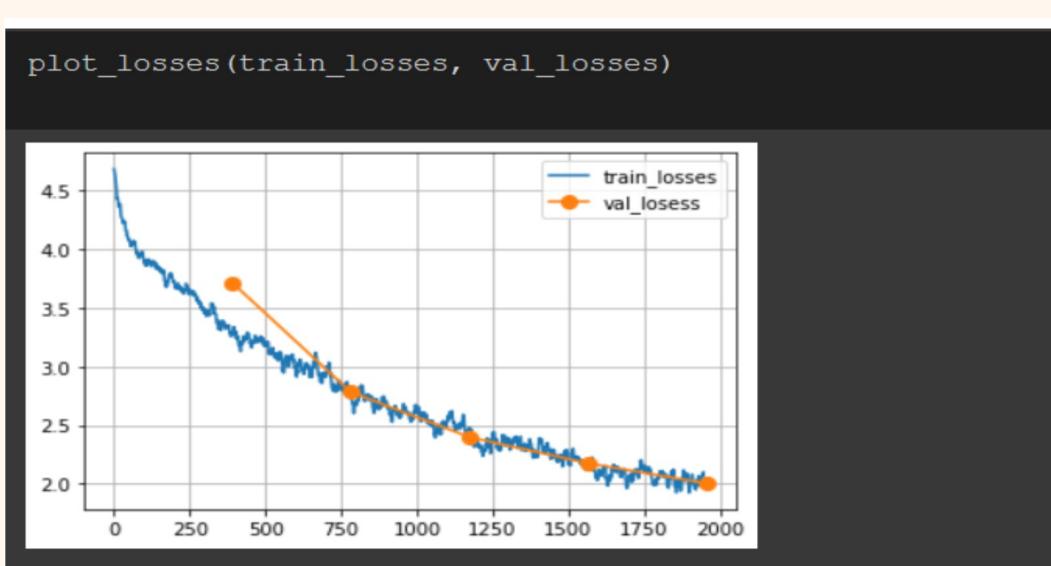
```
start = time.time()
val_acc = get_acc(model, val_loader)
end = time.time()

print('Batch_size={}, epochs={}, lr={}'.format(128, epochs, lr))
print("Val accuracy =", val_acc)
print("Process validation time: {:.4f} s".format(end - start))

HBox(children=[FloatProgress(value=0.0, max=79.0), HTML(value='')])

Batch_size=128, epochs=5, lr=0.001
Val accuracy = 0.4506
Process validation time: 16.6421 s
```

MobileNet2 Validation



```
model = MobileNetV3Small().to(device)
epochs = 5
lr = 1e-3
train_losses, val_losses = train(train_loader, val_loader, model, epochs, lr)
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
EPOCH # 1
train_loss: 3.013 | train_acc: 0.1053
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
val_loss: 3.438 | val_acc: 0.1716
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
EPOCH # 2
train_loss: 2.7053 | train_acc: 0.2164
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
val_loss: 3.089 | val_acc: 0.2381
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
EPOCH # 3
train_loss: 2.7333 | train_acc: 0.2997
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
val_loss: 2.7333 | val_acc: 0.3126
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
EPOCH # 4
train_loss: 2.4096 | train_acc: 0.3647
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
val_loss: 2.4096 | val_acc: 0.3647
HBox(children=[FloatProgress(value=0.0, max=1.0, HTML(value=''))])
MobileNet3 Training
```

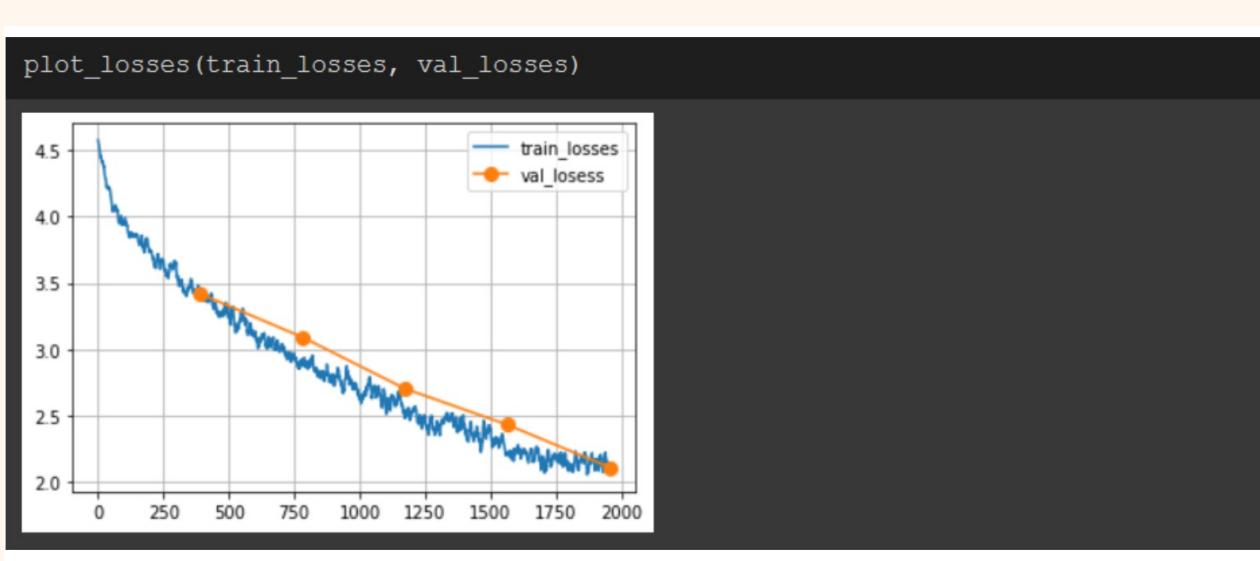
```
start = time.time()
val_acc = get_acc(model, val_loader)
end = time.time()

print('Batch_size={}, epochs={}, lr={}'.format(128, epochs, lr))
print("Val accuracy =", val_acc)
print("Process validation time: {:.4f} s".format(end - start))

HBox(children=[FloatProgress(value=0.0, max=79.0), HTML(value='')])

Batch_size=128, epochs=5, lr=0.001
Val accuracy = 0.4372
Process validation time: 16.0766 s
```

MobileNet3 Validation



```
SAMPLE_IMAGE = 'image1'
IMAGE_URL = 'https://i.stack.imgur.com/CAZAB.png'

def run_visualization(url):
    """Inferences DeepLab model and visualizes result."""
    try:
        f = urllib.request.urlopen(url)
        jpeg_str = f.read()
        original_im = Image.open(BytesIO(jpeg_str))
    except IOError:
        print('Cannot retrieve image. Please check url: ' + url)
        return

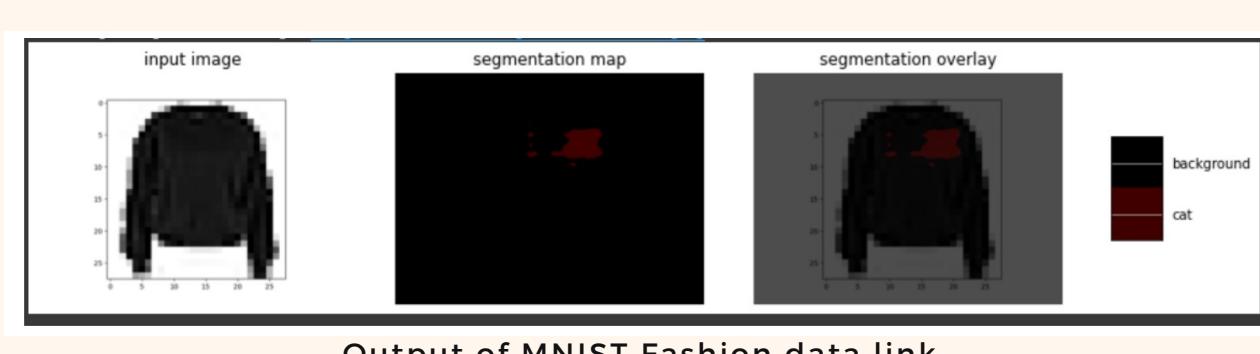
    print('running deeplab on image %s...' % url)
    resized_im, seg_map = MODEL.run(original_im)

    vis_segmentation(resized_im, seg_map)

    image_url = IMAGE_URL or SAMPLE_IMAGE % SAMPLE_IMAGE
    run_visualization(image_url)
```

Test visualization

Segmentation Map and overlayed Output:



Output of MNIST Fashion data link

MobileNetV3 was announced as a new state-of-the-art in mobile classification, detection, and segmentation. Multiple network architectural search techniques, as well as innovations in network design, are being used to produce the next generation of mobile models, as outlined in this study. MobileNetV3 demonstrates how to adapt nonlinearities such as swish and apply squeeze and excite in a quantization-friendly and efficient manner, presenting them as useful tools in the mobile model domain, and I have applied the same nonlinearity in the architecture.

